

Building the SSc-MIM

Construction approach, datasets, and the interaction-validation process

Technical companion to the SSc-MIM overview

Nathan Foulquier — LBAI, UMR 1227 Inserm · CHU de Brest

Detailed reference: [docs/SSc_MIM_construction_and_validation.md](#)

Three governing principles

1. Deterministic curation

Explicit rules (Mazein 2023 + Disease Maps conventions) so two curators produce the same map. HGNC names, fixed compartments, fixed reaction granularity.

2. Never written speculatively

No edge enters the curated map without a verbatim source quote that passes automated gates AND human ratification.

3. Topology is curated, not learned

Mechanisms come from the literature; patient omics ground, validate and read out the map — they never define its structure.

Notation, format & naming conventions

Notation & format

CellDesigner v4.4 · SBGN Process Description glyphs.

SBML Level 2 Version 4 + CellDesigner annotations.

MI2CAST-compliant annotation on every reaction.

Identifiers (HGNC-anchored)

gene/protein → HGNC symbol (TGFB1, SMAD3)

proteoform → symbol+state (SMAD3@S423)

complex → A:B:C (SMAD2:SMAD3:SMAD4)

phenotype → phenotype_<slug>

One reaction per ligand–receptor binding event.

One reaction per kinase–substrate pair (residues collapsed into a state-variable list).


Transcription as TF → mRNA(target), one per target.

Fixed compartment vocabulary (20): extracellular, ECM, plasma membrane, cytosol, nucleus, ER, endosome, mitochondrion, cell-type prefixes...

Modules, crosstalk & phenotype sinks

Figure 1 — SSc-MIM global view (quadrant layout)

526 species · 1488 edges · modules in fixed positions · phenotype sinks centred · inter-module edges shown as curved arcs



M1 Type-I IFN · 19 reactions

M2 TGF-β / fibrosis · 58

M3 EndoMT / vasculopathy · 27

M4 Cytokines IL-6/IL-4/IL-13 · 11

M5 B-cell & autoreactivity · 10

crosstalk · 8

Six phenotype endpoints (disease outputs): myofibroblast activation · ECM/collagen deposition · vascular remodelling · type-I IFN/ISG signature · autoantibody production (autoreactivity) · skin severity (mRSS).

Modelling rule: every entity reaches an endpoint in ≤ 6 steps — verified, 0 violations (make sink-check).

From literature to a released map



Blue = backbone assembly • coral = SSc-specific layer • teal = QC & release.

Each stage is a scripted, reversible step (make targets) with linters and CI — re-running from raw inputs reproduces the map.

The whole automated lane runs end-to-end with `make auto`.

Reactome backbone import & integration

Import canonical pathway scaffolds

M1 type-I IFN · M2 TGF- β (pilot R-HSA-2173789) + PDGF · M3 NOTCH1 (R-HSA-1980143) · M4 IL-6.

Post-process & harmonise

Decode + classify reaction types; rename entities to HGNC; de-duplicate species shared across modules.

Integrate

Merge harmonised module XMLs into one SBML map. Mine embedded PMIDs; auto-fill the bibliography via NCBI E-utils.

reactome_pilot.py	(pilot)
post_process_reactome	(harmonise)
harmonise_imports.py	(harmonise)
integrate_modules.py	(integrate)
extract_pmids_*.py	(pmids)
bib_lookup.py	(bib-lookup)

Backbone = 159 reactions; gives canonical structure, but propagates curator-inference ECO by default.

The SSc-specific curation layer — the novel content

133 hand-curated reactions Reactome cannot structurally hold (autoantibodies, SSc cell states, GWAS-function, clinical crosstalk). Sourced from a systematic 77-paper worklist (24 themes) + co-author full-text PDFs.

Provenance	n	Meaning
original-curation	40	hand-curated (human-original)
fulltext-verified	27	AI-extracted, verbatim-grounded
ai-discovery (gated)	48	via G0-G4 pipeline
claude-reclassify	10	cell-state → conceptual bridge
claude-lit-mine	8	mined, abstract-verified

M2 58 M3 27 M4 21
M1 19 crosstalk 8

Every AI-touched row is tagged, reversible (one TSV row), and carries its grounding quote in notes.

126 / 133 carry a primary PMID.

Provenance & graded evidence

The honest denominator: separate the imported backbone from the genuinely SSc-specific layer.

Layer	Rxns	PMID	Exp. ECO
Reactome backbone	159	99%	0
SSc-specific	133	95%	majority

ECO code	Evidence	n
0000314	direct assay	73
0000033	traceable author stmt	27
0000270	expression pattern	20
0000305	curator inference	9
others (353/315)	interaction / mutant	4

122 confirmed · 6 phenotype_aggregation (definitional sinks) · 4 conceptual_bridge · 1 to_complete (POSTN, under expert review).

make evidence-lint fails CI if any SSc reaction is left untriaged — no undeclared inference debt.

Adding an interaction: the edge-discovery protocol

Governing rule: the curated map is never written speculatively.

- 1 **Stage**
Candidate edge in staging TSV — with a verbatim supporting sentence from the source paper.
- 2 **Fetch**
Cache source text (Europe PMC OA full text / abstract).
- 3 **Gate**
Run G0–G4 automated gates → validation_report.tsv.
- 4 **Ratify**
A human sets decision = promote (default = held).
- 5 **Promote**
Only PASS + promote rows enter the map — tagged, reversible.

Candidates never touch the curated TSV until promoted.

Quality over volume — 12 grounded edges beat 40 shaky ones.

Disease-specific only — content Reactome cannot hold.

One quote → one edge.

Every promoted edge is one removable TSV row.

Five automated gates (G0–G4)

Gate	Check	On failure
G0 schema	required fields present; type in controlled vocabulary	REJECT
G1 HGNC	every gene entity is an official HGNC symbol (HGNC REST)	REJECT (alias→FLAG)
G2 grounding	supporting quote is a verbatim substring of the source text	REJECT — keystone
G3 novelty	gene pair not already an SSc reaction; not a forward Reactome pair	REJECT (dup) / FLAG
G4 evidence	numeric PMID present; corpus text available (SSc context)	FLAG

If the quote is not literally in the paper, the edge is rejected — the core defence against fabricated biology.

Self-test: a fabricated negative-control edge (cand_negctrl, FOXP3→COL1A1, invented quote + decoy PMID) must REJECT on every run; it is hidden from the reviewer app.

Layered QC beyond the gates

Structural

SBML L2V4 validation in CI · sink-connectivity (≤ 6 steps) · CellDesigner-loadability · MINERVA preflight.

Contradiction detector

Flags any gene pair given opposite polarities by different sources → human arbitration. 0 unresolved.

AI verdict pass

Read each interaction's deciding quote + support/contrary abstracts → validate / revise / caution. Now 128 validate · 5 caution.

Citation-integrity sweep

Caught 28 reactions citing an unrelated PMID; every replacement re-verified vs live PubMed. No PMID cited from memory.

Reviewer app

Offline swipe-deck over 143 interactions: verbatim quote, dossier, module map, advisory verdict; decisions export to CSV/JSON.

Human ratification

Binding step: the expert/co-author adjudicates the biology. AI proposes & self-gates; the human decides.

Four public single-cell datasets

197 donors (121 SSc / 76 HC) — skin, lung, blood — re-analysed from raw GEO files. Citations verified against PubMed.

GEO	Reference (verified)	Tissue	Donors (SSc/HC)
GSE195452	Gur C et al. Cell 2022;185:1373	Skin (multiome)	154 (97 / 57)
GSE138669	Tabib T et al. Nat Commun 2021;12:4384	Skin	22 (12 / 10)
GSE128169	Morse C et al. Eur Respir J 2019;54:1802441	Lung (ILD)	13 (8 / 5)
GSE210395	GEO — no linked publication	PBMC	8 (4 / 4)

Skin cohorts ($\approx 100k + 64k$ annotated cells) give the fibroblast-subset resolution that M2/M3 model.

Lung (SSc-ILD) + PBMC extend the map beyond skin. Raw archives SHA-256-pinned in data/MIRROR.md (Zenodo mirror).

How the datasets are used

1

Re-process

scanpy QC, Leiden clustering, cell-type annotation from raw counts.

2

Pseudobulk DEG

Per-donor pseudobulk;
SSc-vs-HC mixed-effects
NB GLM + BH-FDR
($q=0.05$).

3

Map onto species

Match DE genes to map
species → coverage.

4

Score per donor

AUCell → each patient = a
4-module activation
vector.

5

Overlay

60 MINERVA overlays
colour the map by
cluster/dataset.

81.3%

of detectable map species seen in patient data

Per-module coverage: M1 84% · M2 88% · M3 75% · M4 74% · M5 94%.

M1 / type-I IFN significantly up in SSc skin (Gur $p=3.2e-4$; Tabib $p=0.011$).

Cell-type labels validated with CellTypist ($\kappa=0.92$, $ARI=0.70$). Data grounds the map — it does not define its topology.

Per-donor pseudobulk — the right unit of replication

Tens of thousands of cells, but only 4–13 patients per dataset. The biological question is about patients, not cells.

The problem: pseudoreplication

Cells from one patient are not independent (same genotype, biopsy, batch).

Testing per-cell treats them as thousands of samples → inflated power, many false positives (reviewer R2-M1).

The fix

Rebuild one bulk-like profile per patient → the unit of replication becomes the patient, the correct level for SSc-vs-HC.

```
group cells by (donor × cell_type)
row[gene] =  $\Sigma$  raw counts over cells
→ one pseudobulk count row per (donor, cell type).
```

Then DEG on the pseudobulks

NB GLM (pydeseq2) ~ condition

+ Benjamini-Hochberg FDR ($q=0.05$)

Raw counts (not log) — the NB GLM models depth & overdispersion itself.

Trade-off: fewer, more robust hits. Less cell-level granularity, but a valid, reproducible inference — and the same pseudobulks feed AUCell.

AUCell — label-free module activation scoring

“Are a module’s genes concentrated among this patient’s most highly-expressed genes?” If yes → the module is on.

Why — no double-dipping

v1.0 weighted expression by the SSc-vs-HC sign, then tested SSc-vs-HC → circular (reviewer R2-M2).

AUCell scores each donor against the pre-specified module gene set (from the curated map), blind to the SSc/HC label.

Rank-based → robust

Depends only on the order of genes, not absolute values → insensitive to scale, depth, normalisation.

Rank all genes by expression (rank 1 = highest).

Walk the top T (top ~5%); recovery curve $R(r) = \# \text{ module genes with rank } \leq r$.

$AUC = \sum R(r)$, normalised to $[0,1]$

High AUC → module genes cluster at the top → module active.

Output: a 4-module activation vector per patient (+ a Z-score for triangulation).

Feeds the SSc-vs-HC contrasts (M1 IFN $p=3.2e-4$ in skin) and the endotype workflow (per-cluster module fingerprints).

Scripted, tested, containerised, citable

Scripted pipeline

make auto runs the whole lane; make lint / validate run every content + SBML linter.

Tested + CI

pytest suite (no data needed); GitHub Actions: validate_sbml, lint, scripts-smoke, figures.

Containerised

Dockerfile + workflow pin the environment for end-to-end reproduction.

Provenance & metadata

RO-Crate manifest, MIRIAM injection for BioModels, CITATION.cff + .zenodo.json.

Input-data mirror

data/MIRROR.md + .sha256 — SHA-256-pinned GEO archives (Zenodo mirror).

Release

GitHub + Zenodo DOI at git tag v1.0; make release pre-flights.

In one sentence

A literature-curated map, gated against fabrication, grounded in real patient data

Backbone from Reactome → harmonised → integrated into one SBML map (568 sp / 308 rx / 20 comp).

133 SSc-specific reactions added only through G0-G4 gates + human ratification — no edge without a verbatim source.

Four single-cell datasets (197 donors) ground the map: 81.3% coverage, type-I IFN validated in SSc skin.

Fully scripted, tested, containerised and citable (Zenodo DOI, BioModels-ready).