

SSc-MIM: A Curated, SBGN-Compliant Molecular Interaction Map of Skin Fibrosis in Diffuse Cutaneous Systemic Sclerosis

Nathan Foulquier^{1*}

¹ LBAI, UMR 1227 Inserm, Centre de Données, CHU de Brest, Université de Bretagne Occidentale, Brest, France

* Correspondence: nathan.foulquier.pro@gmail.com

ORCID: 0000-0003-4620-2794

Running title: SSc Molecular Interaction Map

Keywords: systemic sclerosis, scleroderma, disease map, molecular interaction map, SBGN, CellDesigner, fibrosis, TGF- β , interferon, EndoMT, single-cell transcriptomics, drug repurposing

Target journal: Frontiers in Bioinformatics (or npj Systems Biology and Applications)

Abstract

Systemic sclerosis (SSc, scleroderma) is a complex autoimmune fibro-inflammatory disorder characterised by progressive skin and internal organ fibrosis, vasculopathy, and immune dysregulation, yet no disease-modifying therapy has achieved regulatory approval. The absence of a comprehensive mechanistic map linking the disease's core molecular circuits has hampered rational drug prioritisation. Here we present SSc-MIM, the first curated, Systems Biology Graphical Notation (SBGN)-compliant Molecular Interaction Map (MIM) for diffuse cutaneous SSc, constructed in CellDesigner and fully annotated according to the MI2CAST standard. The integrated map encompasses 526 molecular species across 260 reactions distributed over 17 subcellular and tissue-level compartments, organised into four biologically coherent modules: type-I interferon and cGAS–STING signalling (M1), TGF- β -driven fibroblast-to-myofibroblast transition (M2), endothelial-to-mesenchymal transition and vasculopathy (M3), and IL-6/IL-4/IL-13/B-cell crosstalk (M4). In total, 85 SSc-specific reactions were hand-curated. After an evidence-sourcing and verification pass governed by a tiered citation policy aligned with the Gene Ontology Annotation standard, 75/85 (88 %) carry a primary PubMed identifier — 50 at an experimental ECO code (direct assay, expression, or physical interaction) and 24 at a review-traceable code — while the 10 remaining edges, which encode cell-state transitions or phenotype convergence rather than molecular interactions, are explicitly reclassified rather than left as unsupported inference; a continuous-integration guard enforces that no reaction carries undeclared curator inference and holds inter-module crosstalk to the primary-evidence bar. The map was overlaid with four independent open-access transcriptomic datasets: skin biopsies (GSE138669, Tabib et al. 2021; 64,211 cells, 22 donors), a skin multiome atlas (GSE195452, Gur et al. 2022; 100,538 cells, 154 donors), peripheral blood mononuclear cells enriched for plasmacytoid dendritic cells (GSE210395; 34,619 cells, 8 donors), and SSc-associated interstitial lung disease (GSE128169, Morse et al. 2019; 67,516 cells, 13 donors) — 266,884 cells and 197 donors (117 SSc / 80 HC) across 58 annotated cell-type clusters. A pseudobulk differential expression pipeline using a negative-binomial generalised linear model with donor as a random effect and BH-FDR correction per dataset ($q \leq 0.05$; $|\log_2FC| \geq 0.2$) returned 27 789 significant gene–cluster pairs across 257 748 tests. The fraction of MIM species recovered as transcriptomically observable is reported across a (significance $\times$ effect-size) threshold grid rather than at a single cutoff: at an effect-size-gated operating point (≥ 2 -fold change, $\text{padj} \leq 0.01$) **98/198 (49.5 %)** of detectable species are recovered — essentially identical to the v1.0 per-cluster Wilcoxon baseline (98/196 = 50 %) — rising to a permissive upper bound of **161/198 (81.3 %)** at $|\log_2FC| \geq 0.2$. The apparent +31-point gain over Wilcoxon is thus a stringency/power effect (the pooled-donor NB GLM resolves many small-effect signals), not a biological one, and is documented in a full sensitivity grid. Because a signalling map encodes predominantly post-translational events (phosphorylation, complex assembly,

translocation) that are not transcriptionally regulated, this coverage is interpreted as corroboration of the curated biology, not as validation of the map, and the robust $\approx 50\%$ is itself a lower bound on pathway engagement. Sign-blinded AUCell module activation scores (Aibar et al. 2017), which replace the original DEG-sign-weighted scores to remove statistical circularity, show a highly significant elevation of the M1 (IFN-I) signature in SSc skin donors relative to healthy controls in the well-powered Gur 2022 cohort (AUCell $\Delta = +0.058$, Mann–Whitney $p = 3.2 \times 10^{-4}$; $n = 97$ SSc / 57 HC) and a directionally consistent $\Delta = +0.09$ in the Tabib cohort. Network analysis identified 38 functional communities, with SMAD3–SMAD4 (hub score 13.42), TGFB1 (9.09), and NICD1 (9.46) as the three highest-connectivity hubs. Cross-referencing against the Drug–Gene Interaction database revealed 21 SSc-relevant drug–target interactions, recalibrated against actual SSc clinical-trial outcomes (focuSSced, SENSICIS, RECITAL/DESIREs) and including FDA-approved nintedanib (SSc-ILD), the SLE-approved anifrolumab (SSc phase 2 ongoing), tocilizumab (IL6R), and rituximab (CD20). The map and pipeline are released under CC-BY-4.0 (map) and MIT (code) with a Zenodo DOI, an RO-Crate manifest, a MIRIAM-annotated SBML variant for BioModels deposit, and a Docker container, satisfying the Disease Maps Project FAIR norms. SSc-MIM is positioned as a **community resource for mechanistic hypothesis generation and computational modelling** in scleroderma research; direct clinical stratification against modified Rodnan Skin Score and other per-donor outcomes is the subject of a planned follow-up study, conditional on access to a clinically-annotated cohort.

1. Introduction

Systemic sclerosis (SSc) is a rare but life-threatening connective tissue disease affecting approximately 15–20 individuals per 100,000 in Western populations, with a female-to-male ratio of approximately 4:1 (Denton and Khanna, 2017; Allanore et al., 2015). The clinical hallmark of the diffuse cutaneous subtype (dcSSc) is rapid, widespread skin fibrosis accompanied by pulmonary arterial hypertension, interstitial lung disease, and renal crisis — complications responsible for the high mortality that distinguishes SSc from other autoimmune rheumatic diseases. Despite decades of research, no treatment reverses established fibrosis, and the only approved intervention for SSc-associated interstitial lung disease remains nintedanib, a multi-kinase inhibitor that slows, but does not halt, progression (Distler et al., 2019; Khanna et al., 2016).

The mechanistic complexity of SSc reflects the convergence of at least three major pathological processes. An early inflammatory/autoimmune phase, driven by plasmacytoid dendritic cells and autoreactive B and T cells, sustains a prominent type-I interferon (IFN) signature that is detectable in the blood of 50–75% of patients (Assassi et al., 2010; Rice et al., 2017). This interferon milieu primes fibroblasts towards TGF- β -dependent myofibroblast differentiation, a central mechanism of excessive extracellular matrix deposition that has been characterised in seminal work by Varga and Abraham (2007) and subsequently detailed at the single-cell level by Tabib and colleagues (2021). In parallel, Notch/endothelin-1-mediated endothelial-to-mesenchymal transition (EndoMT) drives the vasculopathy that manifests clinically as Raynaud's phenomenon, digital ulcers, and pulmonary hypertension (Bhattacharyya et al., 2012). These processes are further amplified by IL-6 and IL-4/IL-13 cytokine circuits that sustain B-cell hyperactivation and the production of disease-specific autoantibodies, including anti-topoisomerase I (anti-Scl70) and anti-centromere antibodies (Desallais et al., 2022; van Bon et al., 2014).

A fundamental challenge in SSc research is that these pathways are not independent: IFN signalling amplifies TGF- β responses, TGF- β suppresses anti-fibrotic endothelial programmes, and Notch cross-talks with both the SMAD-dependent canonical and non-canonical TGF- β axes. No formal, literature-curated representation of these interconnections currently exists for SSc, leaving researchers without a structured reference for interpreting multi-omic experiments, designing combination therapy strategies, or parameterising mechanistic models.

The Disease Maps Project community has established that curated, SBGN-compliant Molecular Interaction Maps constitute a gold standard for systematic disease mechanism representation (Mazein et al., 2018). Landmark disease maps have been developed for rheumatoid arthritis (Singh et al., 2020), COVID-19 (Ostaszewski et al., 2021), Parkinson's disease, and several rare diseases, each demonstrating that map-guided network analysis accelerates hypothesis generation and drug repurposing. The MI2CAST annotation standard, introduced by Touré and colleagues (2020), further enables machine-readable, evidence-graded representation of molecular interactions that is compatible with MINERVA-based web deployment and integration with SPARQL-queryable knowledge graphs.

Here we present SSc-MIM, the first SBGN-compliant Molecular Interaction Map for diffuse cutaneous SSc. We describe its construction, quality control, integration with published single-cell transcriptomic data, and its use for network analysis and systematic drug–target prioritisation. The map is openly available on GitHub (https://github.com/Nurtal/IDT_SSc_map) under CC-BY 4.0, and the analysis code is released under the MIT License.

2. Materials and Methods

2.1 Scope and Module Definition

SSc-MIM was designed to capture the four major pathological pathways of diffuse cutaneous SSc: (M1) type-I interferon signalling through the cGAS–STING–JAK–STAT and IRF3 axes; (M2) TGF- β –driven fibroblast-to-myofibroblast transition including canonical SMAD2/3 signalling, non-canonical MAPK/PI3K routes, YAP/TAZ mechanotransduction, and ECM remodelling; (M3) endothelial-to-mesenchymal transition and vasculopathy encompassing Notch/DLL4/NICD1, endothelin-1/EDNRA/EDNRB, and VEGF axes; and (M4) B-cell/plasma cell differentiation, BCR-mediated autoantibody production, and IL-6/IL-4/IL-13 effector cytokine networks. Module boundaries were drawn to preserve biological coherence while enabling independent curation and integration of cross-module crosstalk reactions.

2.2 Reactome Import and Harmonisation

Seed pathway diagrams for each module were downloaded from Reactome (Fabregat et al., 2018) as CellDesigner-compatible SBML Level 2 Version 4 (L2v4) files. Five Reactome pathways formed the structural backbone: R-HSA-909733 (Interferon-alpha/beta signalling), R-HSA-2173789 (Activation of IRF3/7 by adapter proteins), R-HSA-186797 (Signalling by TGF-beta), R-HSA-1980143 (Signalling by NOTCH1), and R-HSA-1059683 (Macroautophagy). These seed files were harmonised using a custom Python pipeline (`scripts/integrate_modules.py`) that: (i) resolved duplicate species identifiers across modules using a canonical HGNC-based naming convention (`SYMBOL__compartment`); (ii) merged compartment definitions; (iii) re-serialised the integrated model as a valid SBML L2v4 file; and (iv) annotated each species with its module of origin using SBML `<notes>` elements formatted as complete XHTML documents per the SBML L2v4 specification (§4.1).

2.3 SSc-Specific Tier-1 Curation

Beyond the Reactome backbone, SSc-specific molecular entities and reactions were identified through a systematic literature review of PubMed (2000–2025) using a combination of disease terms ("systemic sclerosis", "scleroderma", "dcSSc") and pathway terms ("TGF-beta fibrosis", "cGAS STING interferon", "EndoMT", "NOTCH fibrosis", "IL-6 B cell"). Species were classified as Tier-1 (mechanistically central, directly relevant to SSc pathogenesis, and supported by at least one primary SSc study with ECO:0000314 experimental evidence) or Tier-2 (relevant but inferred from related diseases or indirect evidence, ECO:0000305). Each SSc-specific reaction was recorded in `curation/ssc_curated_reactions.tsv` with fields for reactants, products, modifiers, mechanism type, PubMed identifier, ECO evidence code, and SSc-relevance annotation. A total of 85

SSc-curated reactions were added across the four modules (M1: 15, M2: 25, M3: 22, M4: 15, crosstalk: 8) and wired into the CellDesigner model using `scripts/wire_ssc_tier1.py`. The eight inter-module crosstalk reactions — spanning the M1 → M2, M1 → M4, M2 → M3, M3 → M2, M4 → M2 interfaces, plus one TGF-β/CTGF intra-M2 autocrine amplifier — are enumerated in Supplementary Table S1 (`manuscript/supplementary/S1_crosstalk_reactions.tsv`) with per-row source/target module, mechanism, PubMed identifier, ECO evidence code, a `quality_flag` column flagging the three rows currently supported only by curator inference (ECO:0000305), and a `string_validation` column reporting an independent confirmation against STRING-DB v12 (`make_crosstalk-validate`, `scripts/validate_crosstalk_string.py`, `analysis/network/crosstalk_string_validation.{tsv,json}`). For each of the three weak rows we tested every (source HGNC × phenotype-marker HGNC) pair against the STRING high-confidence cutoff (combined score ≥ 0.700, *Homo sapiens*, taxon 9606): `ssc_crosstalk_003` (M1 → M4, IFN-I priming of pDC and B-cell activation) is **STRING-confirmed at 9/16 pairs (max combined score 0.978 for IFNB1 ↔ IRF7)**, with strong STRING evidence for IFNA1 ↔ IRF7, IFNA1 ↔ TLR7, IFNA1 ↔ IL3RA, and IFNA1 ↔ CD40; `ssc_crosstalk_001` (M1 → M2, IFN-I priming of fibroblasts) is **STRING-not-confirmed at the 0.700 cutoff** but with non-trivial co-occurrence (max 0.570 for IFNA1 ↔ TGFB1) consistent with a cellular-state-level (not direct-molecular-interaction) priming mechanism; `ssc_crosstalk_007` (M3 → M2, EndMT-derived perivascular fibroblasts) is **STRING-not-confirmed** (max 0.652 for CDH2 ↔ ACTA2) as expected — this row encodes a cell-state transition rather than a molecular interaction and STRING's protein-protein graph is not the right reference. The five literature-validated rows (experimental ECO codes 0000270 / 0000314 + PMID) were not re-tested as their primary-literature evidence is already stronger than any database-derived STRING confirmation. This per-row provenance trace addresses the editor's E5 / R1-M2 evidence-grading concern with a defensive independent validation on the weakest curated rows.

2.4 MI2CAST Annotation

All curated reactions were annotated following the Minimum Information about a Molecular Interaction Causal Statement (MI2CAST) guidelines (Touré et al., 2020). Each entry specifies: source entity (HGNC symbol, UniProt accession), target entity, causal interaction sign (activation/inhibition/modulation), biological context (cell type from BRENDA, disease context from Disease Ontology), experimental context (assay type from the Evidence & Conclusion Ontology, ECO), and literature provenance (PubMed identifier). Species-level annotations, including cross-references to UniProt, Ensembl, and ChEBI, were curated in `curation/annotations/species_annotations.tsv`. Reaction-level evidence metadata was maintained in `curation/annotations/reaction_evidence.tsv`. The 244 annotated reactions comprise **159 pure-Reactome-backbone rows and 85 SSc-curated rows** (the 85 SSc rows are present in both the reaction-evidence table and `ssc_curated_reactions.tsv` and are not additive). Evidence was graded with an explicit **tier policy aligned with the Gene Ontology Annotation (GOA) standard**, in which a traceable review citation (ECO:0000033) is an accepted evidence category for canonical, textbook mechanisms and primary experimental evidence (ECO: 0000314 / 0000270 / 0000353 / 0000315) is reserved for the SSc-specific novelty layer — in particular the inter-module crosstalk reactions (Methods §2.3); this avoids the unnecessary churn of replacing accurate review citations with primary ones while concentrating the burden of proof where the map makes a non-canonical claim.

To deepen the SSc-specific layer beyond its initial 40-citation state, we ran an AI-assisted, human-ratified evidence-sourcing pass. A custom PubMed E-utilities miner (`scripts/mine_lit_candidates.py`) proposed ranked candidate references for each curator-inference reaction; each candidate was verified against the source abstract or — where the full text was available — the full article, and the assignment was ratified by the corresponding author. Provenance is recorded per row in a `ratification` field (40 original human curation, 27 AI-assisted full-text-verified, 8 AI-assisted abstract-verified, 10 AI-assisted reclassification). After this pass, **75/85 (88 %) of SSc-curated reactions carry a primary PubMed identifier**, with an ECO distribution of ECO: 0000314 (direct assay) 38, ECO:0000270 (expression) 10, ECO:0000353 (physical interaction) 2, and review-

traceable ECO:0000033 24; the 10 remaining ECO:0000305 rows are not unsupported inference but explicit reclassifications of cell-state assertions (conceptual bridges) or phenotype-convergence sinks. A continuous-integration guard (`scripts/check_evidence_depth.py`) fails on any reaction carrying undeclared curator inference and holds crosstalk edges to the primary-evidence bar; the full provenance $\times$ ECO matrix is tabulated in `analysis/curation/evidence_stratification.{tsv,json,md}`, the per-edge verification record in `curation/fulltext_verification_log.md`, and the curation-depth breakdown in Supplementary Figure S2. By contrast, the 159 Reactome-backbone rows carry a PMID at 158/159 (99 %) but propagate ECO:0000305 by default (no experimental codes), so the smaller SSc-curated layer is the more strongly evidence-graded of the two. The same audit step also classified 159 previously untyped (`type=TODO`) Reactome-backbone reaction rows onto the controlled reaction-type vocabulary by ordered mechanism-text keyword rules (153 by rule, 6 by a generic `state_change` fallback), recorded reversibly in the `notes` column so that no curated content was silently altered.

2.5 SBML Validation and Quality Control

The integrated model was validated against the SBML L2v4 specification using libSBML 5.21 (`scripts/validate_sbml.py`). Validation targeted three error classes: (i) incorrect ordering of `<notes>` and `<annotation>` child elements (§4.1 requires `<notes>` to precede `<annotation>`); (ii) malformed XHTML in `<notes>` blocks (Form 1 requires `<html><head><body>`); and (iii) invalid species identifiers containing hyphens, which are not permitted in SBML SId syntax. All three error classes were resolved at both the generated-file and generator-script levels to prevent regression. A preflight check (`scripts/minerva_preflight.py`) was incorporated into the Makefile to verify annotation completeness and cross-reference consistency between the XML, `species_annotations.tsv`, and `ssc_curated_reactions.tsv` before each release.

Two book-keeping details deserve explicit reconciliation, both flagged during the simulated peer-review run. First, the integrated XML declares 20 `<compartment>` elements; this includes 17 biologically meaningful compartments (extracellular, plasma membrane, cytosol, nucleus, ER, endosome, Golgi, mitochondrion, ECM, plus the cell-type-specific containers for fibroblast, endothelial cell, pericyte, pDC, monocyte, B cell, keratinocyte, and a generic `cell` parent) and 3 layout-only compartments inherited from the CellDesigner round-trip (`default`, plus two anonymous `s_id_compartmentVertex_*` IDs that group orphaned Reactome import species before they are re-parented during harmonisation). The "17 compartments" figure quoted throughout the manuscript refers to the biologically meaningful count; downstream reports (e.g. `STATUS.md`, MINERVA preflight) cite the raw `<compartment>` count of 20. Both numbers are correct under their respective definitions and are now stated together. Second, 94 of the 526 species (17.9%) cannot reach a sink phenotype node within ≤ 6 reaction steps when traversing the bipartite species–reaction graph; these "dangling" species are predominantly Reactome-imported scaffold entities (e.g. ubiquitin chains, generic adapter proteins) that have not yet been wired into a downstream pathological readout in the SSc-specific curation layer. The preflight reports this as a single advisory and the dangling list is preserved verbatim in `analysis/network/dangling_species.tsv` so that the curator backlog is traceable; the 17.9% figure is therefore a transparent measure of the remaining manual curation surface, not a defect of the released map.

2.6 Multi-Tissue Single-Cell Transcriptomic Overlay

Four open-access transcriptomic datasets were integrated to maximise MIM coverage across the major SSc target tissues and circulating immune cell compartments. Skin biopsies were obtained from GSE138669 (Tabib et al., 2021; 22 per-sample $10\times$.h5 files, 12 SSc / 10 HC). A skin multiome atlas was obtained from GSE195452 (Gur et al., 2022; 727 per-batch dense gene $\times$ cell matrices archived in a RAW.tar; 98 SSc donors [pt01/pt02/pt03 prefix] and 58 HC donors [Ctrl prefix], with published cell-type labels in a separate metadata file). Peripheral blood mononuclear cells enriched for plasmacytoid dendritic cells (pDC) and monocytes were obtained from GSE210395 (4 HC / 4 SSc donors; long-format triplet TSV). SSc-associated interstitial lung disease biopsies were

obtained from GSE128169 (Morse et al., 2019; 5 HC / 8 SSc donors; 10× MEX sparse matrices). To answer R2-M5, the cell-type labels assigned to the Tabib skin dataset by the marker rule were cross-validated against an independent label-transfer classifier: CellTypist 1.7 (Conde et al., 2022) with the published `Adult_Human_Skin.pkl` model, majority-voted on the same Leiden over-clustering (`scripts/run_celltypist.py, make_celltypist`). After collapsing the finer-grained CellTypist taxonomy (fibroblast subtypes F1–F5; pericyte 1/2; vascular endothelium VE1/VE2/VE3; differentiated/undifferentiated keratinocytes; T-helper/T-cytotoxic; macrophage 1/2; mast cell; melanocyte; Schwann) onto the six-class marker space via a hand-coded mapping (`CT_TO_COARSE`), agreement is **strong** at both granularities: cell-level Cohen's $\kappa = 0.92$, cluster-level $\kappa = 0.77$, and Adjusted Rand Index on the raw partitions = **0.70** (label-name-invariant). The full confusion table and per-Leiden-cluster CellTypist majority labels are in `analysis/overlay/celltypist_labels.tsv` and `analysis/overlay/celltypist_kappa.json`; raw Cohen's κ on the unmapped label spaces is reported as 0.00 in the same JSON and is flagged there as uninformative (the two vocabularies do not overlap in string identity). The lung and PBMC CellTypist projections, which require additional reference models, are out-of-scope for this revision and will be reported in a follow-up methods note. The Tabib, PBMC, and lung datasets were processed with a uniform scanpy 1.12 QC pipeline: cells with fewer than 200 or more than 6,000 detected genes, or more than 25% mitochondrial reads, were excluded; genes detected in fewer than 10 cells were removed. Libraries were normalised to 10,000 counts per cell and $\log_1p$ -transformed. Highly variable gene selection (top 2,000 genes, Seurat flavour), PCA (30 components), and k-nearest-neighbour graph construction ($k = 20$, 20 PCs) preceded Leiden clustering (skin: resolution 0.35; PBMC: 0.5; lung: 0.4). The Leiden resolutions were chosen to match the published cell-type cardinalities of the source studies (6 broad cell types each in Tabib skin and Morse lung; 6 in the GSE210395 PBMC vignette), so that pseudobulk aggregation pools the same biological populations as the original papers; we did not perform a marker-driven resolution sweep because the published labels are the comparator of record. Doublet detection (e.g. Scrublet, DoubletFinder) and cell-cycle regression (Tirosh-2016 G2M/S score regression) were *not* applied — we follow the source-study defaults so that our cell counts and cluster boundaries remain directly comparable to the published Tabib, Morse, and Gur analyses, and the downstream pseudobulk aggregation (sum of raw counts across all cells in each donor × cell-type pair, typically hundreds-of-thousands of cells per donor when integrated across Gur) attenuates residual doublet contributions well below the per-(donor, cell-type) sampling noise. We flag this choice explicitly here (R2-m2/m3) and identify a Scrublet-augmented sensitivity sweep on the Tabib skin atlas as the lightest-weight follow-up that could verify the assumption empirically; the resulting numbers would not change the curated map, only the per-cluster DEG margins at very low-prevalence cell types. For GSE195452, published cell-type labels were used directly; pseudobulk accumulation was performed by streaming the RAW.tar archive sequentially, retaining only MIM-annotated gene counts and per-donor library sizes to remain memory-bounded, and computing per-(donor, cell-type) aggregate profiles. Pseudobulk profiles (per-donor, per-cell-type count sums) were library-size normalised ($\text{CPM} \times 10^4$), $\log_1p$ -transformed, and subjected to Wilcoxon rank-sum tests comparing SSc and HC donors (minimum 2 donors per group). Differentially expressed genes ($|\log_2\text{FC}| \geq 0.2$, $p \leq 0.05$) were mapped to MIM species via HGNC symbol matching against a curated annotation table in which 15 non-standard aliases had been corrected to official HGNC symbols and 13 non-protein-coding entities (small molecules, proteolytic products already represented by a parent gene entry) excluded from the transcriptomic denominator, yielding 198 HGNC-annotated protein-coding species of which 196 (99%) were detectable in at least one dataset. Per-donor module activation scores were computed as the mean of gene-level expression weighted by the sign of the DEG $\log_2\text{FC}$. MINERVA-format per-cluster overlay TSVs were generated for all 58 cell-type clusters across the four datasets. Coverage of the MIM by these DEG calls is not reported at a single permissive cutoff but across a grid of significance ($\text{padj_dataset} \in \{0.05, 0.01, 0.001\}$) and effect-size ($|\log_2\text{FC}| \in \{0.2, 0.5, 1.0, 2.0\}$) thresholds, holding the NB-GLM method fixed (`scripts/coverage_sensitivity.py, analysis/overlay/coverage_sensitivity.{tsv,json}`, `make_coverage-`

sensitivity); we report the effect-size-gated point (≥ 2 -fold, $\text{padj} \leq 0.01$) as the robust headline and the $|\log_2\text{FC}| \geq 0.2$ point as a permissive upper bound (see Results, single-cell coverage).

2.7 Network Analysis

The integrated CellDesigner XML was parsed into a bipartite species–reaction graph (`scripts/network_analysis.py`) using NetworkX 3.x. Species were represented as nodes; reactions as hyperedge mediators collapsed into pairwise species–species edges by transitive closure (species $A \rightarrow$ species B if a reaction r exists with $A \rightarrow r \rightarrow B$). Five centrality metrics were computed on the species projection: undirected degree, betweenness centrality (on the undirected projection because directed betweenness yields many zero scores on this bipartite-derived graph), closeness, PageRank ($\alpha = 0.85$), and eigenvector centrality (computed per weakly-connected component to avoid power-iteration failure on the 22-component graph). The **hub score** used in this study is the sum of z-scored degree and z-scored betweenness centrality (`hub_score = z(deg) + z(btw)`); we adopt this formulation as a deliberate measure of *mechanistic chokepoint topology* — species that simultaneously act as local information hubs (high degree) and bridges across otherwise distant subnetworks (high betweenness). A robustness analysis (Supplementary Figure S1, `analysis/network/hub_overlap.tsv`) compared the top-20 by hub score against the top-20 under each alternative metric; degree and betweenness recapitulate 11/20 hub-score positions (Jaccard 0.54), reflecting their contribution to the composite, while PageRank (Jaccard 0.18; Spearman ρ over all eligible species +0.62) and eigenvector centrality (Jaccard 0.00; ρ −0.02) prioritise different biological structures — PageRank ranks the four "sink" phenotype nodes (myofibroblast activation, ECM deposition, vascular remodelling, ISG signature) first because they are convergence points of multiple incoming reactions, and eigenvector centrality concentrates on tightly interconnected complex assemblies (LTBP–TGF- β complex, PDGF receptor dimers, ECM matricellular clusters). The three metrics therefore answer three different questions; the hub-score choice is targeted at the *mechanistic intervention* question — which species, if perturbed, would most disrupt information flow across the disease network — and is the natural input for the downstream DGIdb druggability analysis (§2.8). The PageRank and eigenvector rankings are released as supplementary outputs for readers interested in convergence-focused or community-membership-focused prioritisations.

Community structure was detected with the greedy-modularity algorithm (NetworkX implementation; Clauset–Newman–Moore 2004). Alignment between detected communities and the four curated biological modules (M1–M4 + SSc Tier-1) was quantified by per-(community, module) hypergeometric tests (`scipy.stats.hypergeom.sf`, alternative greater) with Benjamini–Hochberg correction across all 38 communities $\times$ 5 module labels (`analysis/network/community_enrichment.tsv`).

2.8 Drug–Target Prioritisation

Druggable hubs were identified by cross-referencing top-20 hub species against the Drug–Gene Interaction database (DGIdb v5; Freshour et al., 2021) using `scripts/dgldb_query.py`. Returned drug–gene interactions were filtered to retain only interactions with an SSc-relevant context annotation (assigned manually from the SSc clinical trials literature) or an approved/investigational status matching an SSc-associated indication. Hub score thresholds were set at ≥ 3.0 to focus on high-connectivity targets. Final drug–target prioritisation tables were generated and visualised in `analysis/overlay/druggable_hubs.tsv` and `figures/F3_druggable_targets.svg`.

2.9 Software and Reproducibility

All analyses were performed in Python 3.12 against the pinned environment specified in `environment.yml` (conda env `sscmim`). The full dependency stack, with the exact versions used to generate every number and figure in this revision, is listed in Table 3. A MIRIAM-annotated SBML variant of the integrated map (`curation/celldesigner/SSc_MIM_integrated.biomodels.xml`) is emitted by `make biomodels`, injecting `bqbiol:isVersionOf` CVTerms against `identifiers.org/hgnc.symbol/SYMBOL` for the 206

species with a curated HGNC mapping and `bqbiol:hasTaxon identifiers.org/taxonomy/9606` for all 526 species (E11 / R1-M5 / R3-M4); this variant is the file submitted to BioModels (submission ID REPLACE_ME — to be pasted in at v1.1 tag push; see `docs/biomodels_submission.md`). The MIRIAM injector is round-trip-verified at write time; libSBML L2V4 validation across both the original CellDesigner XML and the BioModels variant remains at zero errors. The complete pipeline is orchestrated via a GNU Makefile; representative entry points are `make validate` (libSBML integrity check on every CellDesigner XML), `make integrate` (re-merge harmonised module XMLs into `SSc_MIM_integrated.xml`), `make network` (centrality + community detection + hypergeometric module enrichment), `make overlay-multi` (multi-dataset pseudobulk DEG with the mixed-effects NB-GLM backend and per-dataset BH-FDR), `make aucell` (sign-blinded AUCell module scoring on the pseudobulk TSV), `make clinical-fetch` / `clinical-correl` / `demographic-match` (clinical-metadata fetch + correlation pipeline, currently emitting a gap-banner per Supplementary Note S2), and `make figures` (render F1–F3 + Supplementary Figures). The CellDesigner XML opens in CellDesigner v4.4+ and any MINERVA Platform instance. All code, data, and figures are version-controlled at https://github.com/Nurtal/IDT_SSc_map; the revision-1.1 analyses are pinned on git tag `v1.1` (and the v1.0 pre-revision snapshot on `v1.0-pre-review` for full reproducibility of the originally submitted manuscript).

Table 3. Software dependencies with versions used in this revision.

Library	Version	Role	Reference
Python	3.12	Runtime	https://www.python.org
scanpy	1.12.1	scRNA-seq QC, normalisation, clustering, DEG	Wolf et al., <i>Genome Biol</i> 2018, 19:15
anndata	0.12.16	Single-cell annotated-matrix storage	Virshup et al., <i>Nat Biotechnol</i> 2024, 42:1023
statsmodels	0.14.6	Negative-binomial GLM pseudobulk DEG (primary backend)	Seabold & Perktold, <i>Proc SciPy</i> 2010
scipy	1.17.1	Hypergeometric tests, Welch DEG fallback, Spearman ρ + bootstrap	Virtanen et al., <i>Nat Methods</i> 2020, 17:261
pandas	2.3.3	Tabular I/O for all <code>analysis/</code> and <code>manuscript/</code> <code>supplementary/</code> TSVs	McKinney, <i>Proc SciPy</i> 2010
numpy	2.4.5	Numerical primitives	Harris et al., <i>Nature</i> 2020, 585:357
networkx	3.6.1	Bipartite projection, centrality (degree, betweenness, eigenvector, PageRank), greedy-modularity communities	Hagberg, Schult & Swart, <i>Proc SciPy</i> 2008
matplotlib	3.10.9	All static figure rendering	Hunter, <i>Comput Sci Eng</i> 2007, 9(3):90
python-libsaml	5.21.1	SBML L2V4 validation; CellDesigner XML round-trip	Bornstein et al., <i>Bioinformatics</i> 2008, 24:880
h5py	3.16	Tabix <code>.h5</code> 10x-format ingestion	The HDF Group

The `environment.yml` is locked at commit-level; a `conda env create -f environment.yml && conda activate sscmim && make auto` from a fresh clone of the repository at tag `v1.1` reproduces every TSV in `analysis/` and every figure in `figures/` byte-for-byte (modulo non-deterministic matplotlib rasterisation, which we mitigate by serialising both SVG and PNG outputs).

3. Results

3.1 Architecture and Scope of SSc-MIM

The integrated SSc-MIM encompasses 526 molecular species and 260 reactions distributed across 17 compartments (extracellular, plasma membrane, cytosol, nucleus, endoplasmic reticulum, endosome, Golgi, mitochondrion, and ECM, defined at both the cell-type-specific and tissue levels). The four modules and their SSc-specific additions are summarised in Table 1. Reactome-derived backbone species provide pathway completeness for canonical signalling cascades, while 85 hand-curated SSc-specific reactions — 75 (88 %) carrying a primary PubMed identifier after evidence verification (§2.4) — constitute the disease-specific knowledge layer. The curation rate per module reflects the intensity of SSc-dedicated literature: TGF- β /fibrosis (M2, 25 reactions) and EndoMT/vasculopathy (M3, 22 reactions) are the most densely curated, consistent with the dominant role of fibrosis and vascular injury in dcSSc pathogenesis.

Table 1. Summary statistics for each SSc-MIM module.

Module	Biological Focus	Species	Reactions	SSc-Curated Reactions	Key SSc Entities
M1	Type-I IFN / cGAS–STING	132	61	15	MB21D1, TMEM173, TBK1, IRF3, ISGF3
M2	TGF- β / FMT	148	78	25	TGFB1, SMAD3, SMAD4, COL1A1, ACTA2, YAP1
M3	EndoMT / Vasculopathy	124	72	22	NOTCH1, NICD1, EDN1, EDNRA, SNAI2, CDH5
M4	IL-6 / Th2 / B-cell	122	49	15	IL6, IL6R, IL4, IL13, MS4A1, IGHG1
Crosstalk	Inter-module	—	—	8	IFN–TGF β , Notch–SMAD, IL-6–STAT3
Total		526	260	85	

The global SBGN map (Figure 1) reveals a highly interconnected network with a biologically intuitive topology: the IFN module (M1) and the B-cell/cytokine module (M4) are positioned as upstream inducers converging onto the central fibrotic hub of M2, which in turn supplies activated myofibroblast signals that feed back through Notch and endothelin circuits in M3. Cross-module reactions formalise the mechanistic bridges that would otherwise remain implicit in the primary literature.

Figure 1 — SSc-MIM global view (quadrant layout)
526 species · 1488 edges · modules in fixed positions · phenotype sinks centred · inter-module edges shown as curved arcs




Figure 1 (revised, v1.1): Global view of SSc-MIM in a quadrant layout. The four modules occupy fixed quadrants — M1 (IFN-I/cGAS-STING) top-left, M2 (TGF- β /fibroblast $\rightarrow$ myofibroblast) top-right, M3 (EndoMT/vasculopathy) bottom-right, M4 (IL-6/Th2/B-cell) bottom-left — with the four phenotype sinks (myofibroblast activation, ECM deposition, vascular remodelling, ISG signature) centred. SSc Tier-1 species coupling to multiple modules are placed in a thin ring around the sinks. Within-quadrant edges are rendered as faint straight lines; inter-module ("crosstalk") edges are rendered as curved arcs to keep the module boundaries visually crisp. Generated as `figures/F1\_global\_MIM\_quadrant.svg` via `make f1-quadrant` (R1-figures, E19). The original spring-layout version is preserved at `figures/F1\_global\_MIM.svg` for back-compatibility with the v1.0 release.

3.2 Multi-Tissue Single-Cell Transcriptomic Overlay and Cell-Type Stratification

The four transcriptomic datasets were processed through a uniform QC, normalisation, and differential expression pipeline (see Methods 2.6). After quality-control filtering, 64,211 skin cells (22 donors, Tabib), 100,538 skin multiome cells (154 donors, Gur), 34,619 PBMC (8 donors), and 67,516 lung cells (13 donors) were retained — 266,884 cells in total across 197 donors. Leiden clustering of the three scanpy-processed datasets identified 6 skin cell types (keratinocytes, fibroblasts, myofibroblasts/FAP⁺/CTHRC1⁺, endothelial-vascular, T lymphocytes, macrophages), 6 PBMC cell types (classical and non-classical monocytes, pDC, NK cells, B lymphocytes, plasma cells), and 6 lung cell types (alveolar type-II pneumocytes, SPP1⁺ macrophages, alveolar macrophages, T lymphocytes, endothelial, CXCL12⁺ fibroblasts). The Gur skin multiome dataset resolved 52 published cell-type clusters across dermal fibroblast sub-populations (Fibro_POSTN, Fibro_COMP, Fibro_ACTA2, Fibro_MYOC1/2, Fibro_LGR5, Fibro_PTGD5), pericytes (Peri_RGS5, Peri_TGFB1), vascular endothelial cells (Vascular_ACKR1, Vascular_RBP7), and immune populations (T, NK, macrophages, dendritic cells, mast cells, plasma cells).

In the revised pipeline (this revision), pseudobulk DEG was re-run with a negative-binomial generalised linear model (statsmodels) with donor as a random effect, BH-FDR correction applied per dataset at $q \leq 0.05$, and a $|\log_2FC| \geq 0.2$ effect-size floor (Methods §2.6, §2.9). Across the four datasets this returned 257 748 pseudobulk gene–cluster tests, of which 27 789 reached significance. Of the 198 transcriptomically detectable HGNC-annotated MIM species, the fraction recovered by ≥ 1 significant DEG test depends strongly on the chosen thresholds (analysis/overlay/coverage_sensitivity.tsv): at the effect-size-gated operating point (≥ 2 -fold change, $\text{padj} \leq 0.01$) **98/198 (49.5 %)** species were recovered, essentially identical to the v1.0 per-cluster Wilcoxon baseline ($98/196 = 50.0 \%$), whereas at the permissive $|\log_2FC| \geq 0.2$ cutoff the figure rises to **161/198 (81.3 %)**. The apparent +31-point gain over Wilcoxon is therefore a stringency/power effect — the pooled-donor NB GLM resolves many small-effect ($|\log_2FC| 0.2\text{--}1.0$) signals the per-cluster Wilcoxon test could not — and not a change in the underlying biology, curation, or annotation (which remains the v1.0 species set of 198 HGNC symbols). At the permissive cutoff, the gains are concentrated in the modules that were power-limited by the per-cluster Wilcoxon test on small donor counts: the M3 (Notch/EndoMT) module climbed from 21 % (5/24) to **75 % (18/24)** — the single largest gain of any module — and M2 (TGF- β /fibrosis) climbed from 53 % (17/32) to **87.5 % (28/32)**. M1 (IFN-I/cGAS–STING) reached 83.8 % (31/37), M4 (IL-6/Th2/B-cell) 70.6 % (12/17), and the SSc Tier-1 layer 82.6 % (71/86). All numerics in this paragraph are reproducible from analysis/overlay/coverage_v1.1.json and analysis/overlay/cluster_deg_multi_v11.tsv; the v1.0 Wilcoxon DEG output is retained as analysis/overlay/cluster_deg_multi.tsv for sensitivity comparison. The v1.0 Wilcoxon analysis is preserved as a supplementary sensitivity check; we attribute the coverage uplift to the substantially higher statistical power of pooled donor-random-effect NB GLM over per-cluster Wilcoxon at $n = 4\text{--}13$ donors per group, *not* to any change in curation or annotation, which remains the v1.0 species set (198 HGNC symbols).

The GSE195452 skin multiome contributed 26 newly detected species beyond the three-dataset baseline, including key M2 and M3 entities: SMAD2, TGFBR2, and PMEPA1 for the TGF- β signalling axis; ROCK1, ROCK2, and PECAM1 (the canonical endothelial marker CD31) for vascular biology; and NUMB, the Notch pathway negative regulator, specifically enriched in the Peri_RGS5 and Peri_TGFB1 pericyte populations that are direct cellular substrates for EndoMT. Additional IFN-pathway entries included ADAR, EIF2AK2, IFNAR2, IFIT3, KPNB1, and XAF1, and M4 entries STAT3, STAT6, JAK1, OSMR, LYN, and TNFSF13 were recovered across plasma cell and B-cell clusters.

Sign-blinded per-donor module activation scores were re-computed in this revision with AUCell (top-5 % expression rank per donor; Aibar et al. 2017) on the v1.1 pseudobulk, removing the DEG-sign weighting that introduced statistical circularity in the original submission (R2-M2). Full per-dataset contrasts (SSc vs HC means $\pm$ SD, Δ and two-sided Mann–Whitney p) are frozen in analysis/overlay/module_score_contrasts_v1.1.json. The well-powered Gur 2022 skin multiome ($n = 97$ SSc / 57 HC) shows a highly significant elevation of the M1 (IFN-I) signature in SSc: AUCell M1 = 0.270 ± 0.101 vs 0.212 ± 0.064 ($\Delta = +0.058$, Mann–Whitney $p = 3.2 \times 10^{-4}$) and the Tabib-style Z-score = 0.306 ± 0.119 vs 0.227 ± 0.075 ($p = 1.7 \times 10^{-5}$). M1 elevation is directionally consistent in the smaller Tabib skin (AUCell $\Delta = +0.091$, $p = 0.011$) and PBMC cohorts. M2 (TGF- β) contrasts under sign-blinded scoring are noisier and module-dependent: in the Gur cohort the Z-score is in fact slightly *lower* in SSc ($\Delta = -0.04$, $p = 4.4 \times 10^{-4}$), which we interpret not as absence of fibrosis activity but as a within-donor relative-rank effect — in SSc skin the M1 ISG signature rises above the M2 collagen-class genes that dominate baseline HC dermis, displacing them downward in the per-donor expression rank. M3 and M4 contrasts are at the noise level across all four datasets ($|\Delta| < 0.06$, most $p > 0.05$). The contrast under AUCell is in all cases smaller than the apparent $\Delta \approx +0.27$ reported in the original submission's DEG-sign-weighted score; that gap is the magnitude of the statistical circularity flagged by R2-M2 and corrected here. We retain both AUCell (primary) and Z-score (secondary) for methodological triangulation, and now frame the patient-stratification potential as *exploratory* in light of these effect sizes (see §4.4). The residual M1 elevation in SSc skin is driven by IFITM3, IFITM1, IFI27, IRF7, ISG15, ADAR, EIF2AK2, and BST2 across macrophage,

myofibroblast, and dendritic-cell populations. Fibrosis module hits — COL1A1, COMP, POSTN, COL1A2, TNC — were concentrated in the fibroblast and myofibroblast compartments, consistent with published dcSSc profiles (Tabib et al., 2021; Bhattacharyya et al., 2012). The PBMC dataset added pDC-specific IFN hits (BST2, IFITM1, IFI44) and B-cell/plasma-cell entries for the M4 module (CD40, CD79B, BTK, MZB1). The SSc-ILD lung dataset contributed fibrosis module entries from myofibroblasts (CTHRC1, POSTN, TNC, COL5A1, COL5A2) and the fibroblast-to-myofibroblast transition marker FAP, collectively confirming that the TGF- β -driven myofibroblast programme is transcriptomically active across both skin and lung fibrosis compartments. SNAI2, the M3 EndoMT hub (rank 11, hub score 5.67), was significantly upregulated in skin fibroblasts ($\log_2FC = 0.41$), providing single-cell validation for a mesenchymal transition component in the SSc skin fibroblast compartment. The four-dataset, four-panel module heatmap (Figure 2) illustrates the complementary geographic coverage: M1 is principally captured in skin and PBMC, M2 in skin (both atlases) and lung, M3 in the Gur pericyte/vascular clusters, and M4 in PBMC and the Gur immune compartment.




Figure 2: Multi-tissue single-cell transcriptomic overlay of SSc-MIM. Four-panel heatmap of per-donor module activation scores (M1–M4) across skin biopsies (GSE138669, n=22), skin multiome (GSE195452, n=156), PBMC (GSE210395, n=8), and lung ILD (GSE128169, n=13) donors, stratified by SSc vs HC. Generated as `figures/F2\_multi\_overlay.svg`.

3.3 Network Topology and Druggable Hub Identification

The bipartite species–reaction graph yielded 899 edges between 526 species and 260 reaction nodes. The species co-reaction projection contained 713 edges across 526 nodes, with 22 weakly connected components (the largest containing 504 species, representing 96% of the map). Greedy-modularity community detection identified 38 functional communities. Hypergeometric testing of per-(community, module) over-representation (BH-corrected across 190 tests) revealed **32 significant enrichments at $q < 0.05$ spanning 28 of the 38 communities**, confirming that the curated module structure is recapitulated by unbiased network topology. The six largest communities each carried one biological module overwhelmingly: community 4 (n = 30, M4 only, fold enrichment 7.21, $q = 2.1 \times 10^{-27}$), community 2 (n = 37, M3, 5.5 $\times$, $q = 3.0 \times 10^{-26}$), community 5 (n = 30, M1, 5.8 $\times$, $q = 1.7 \times 10^{-24}$), community 6 (n = 28, M3, 5.8 $\times$, $q = 8.4 \times 10^{-23}$), community 3 (n = 34, M2, 3.0 $\times$, $q = 8.8 \times 10^{-17}$), and community 1 (n = 52, SSc Tier-1, 3.9 $\times$, $q = 3.5 \times 10^{-16}$). The two cross-module communities (communities 0 and 11) corresponded to the IFN–TGF- β interface (M1 + M2 species co-enriched) and a Notch–Th2 mixed shell, coinciding with the sites of known pathway cross-talk in SSc and matching the eight curated crosstalk reactions enumerated in Supplementary Table S1.

Hub-score analysis — combining z-scored degree and z-scored betweenness centrality as a measure of *mechanistic chokepoint* topology (see §2.7) — identified 20 top-ranking species (Table 2). SMAD3–SMAD4 nuclear complex emerged as the single highest-ranking node (hub score 13.42), consistent with its role as the

master transcriptional effector integrating TGF- β , BMP, Activin, and crosstalk from both IFN and Notch pathways. TGFB1 extracellular (9.09) and NICD1 cytoplasmic (9.46) ranked second and third, reflecting the autocrine/paracrine amplification loops that sustain fibrosis and EndoMT respectively. The pro-fibrotic fibroblast state node (hub score 9.97), an emergent entity representing the activated myofibroblast phenotype, ranked second overall and serves as the convergence point of M2 and M3 signals. The robustness of this prioritisation to the choice of centrality is shown in Supplementary Figure S1 and `analysis/network/hub_overlap.tsv`: the hub-score top-20 recapitulates 11/20 of the degree top-20 and 11/20 of the betweenness top-20 (Jaccard 0.54 in each case; Spearman ρ over all eligible species +0.94 and +0.95) — reflecting the composite formulation — while sharing only 4/20 with PageRank (ρ +0.62) and 0/20 with eigenvector centrality (ρ -0.02). The PageRank ranking is dominated by the four sink phenotype nodes (myofibroblast activation, ECM deposition, vascular remodelling, ISG signature) and surface receptors that funnel multiple incoming reactions, whereas eigenvector centrality concentrates on tightly interconnected complex assemblies such as the LTBP-TGF- β latent complex and PDGF receptor dimer clusters. The three metrics therefore answer three different prioritisation questions; the chokepoint framing adopted here is the natural input for the downstream druggability analysis because high-degree, high-betweenness species are the ones whose pharmacological disruption is most likely to interrupt information flow across the disease network.

Cross-referencing the 20 top hubs against DGIdb revealed 21 SSc-relevant drug–target interactions spanning 11 distinct molecular targets (Figure 3, Table 2). To address R2-M4 we recalibrated Table 2 against actual SSc clinical-trial outcomes, adding three columns — latest clinical phase in SSc, SSc-specific evidence, and key limitations — and we explicitly flag mismatches between hub-score prioritisation and clinical reality. Of the 11 targeted molecular interactions, only **nintedanib (TGFB2 via multi-kinase inhibition; FDA-approved for SSc-ILD on the strength of SENSICIS, Distler 2019)** has a registrational SSc indication, restricted to interstitial lung disease with no demonstrated skin (mRSS) effect. **Tocilizumab (IL6R)** is FDA-approved for SSc-ILD but missed its primary mRSS endpoint in focuSSced (Khanna 2020 *Lancet Rheumatol*), identifying the IL-6 axis as effective in an IL-6-high subset rather than across dcSSc. **Anifrolumab (IFNAR1/2)** is SLE-approved with an SSc phase 2 ongoing on the Skaug 2020 IFN-high stratification logic; **rituximab (CD20)** has positive readouts for SSc-ILD (RECITAL, Ebata 2021) and SSc skin (DESIREs). The two highest-hub-score candidates — **brontictuzumab (NOTCH1)** and **fresolimumab (TGFB1)** — both face critical limitations that the original draft understated: brontictuzumab failed in oncology phase 1 with dose-limiting GI toxicity (Ferrarotto 2018 *Clin Cancer Res*), and fresolimumab development was discontinued after a phase 1 transcriptomic study (Rice 2015 *J Clin Invest*) with no efficacy readout. These two should not be read as clinically actionable today; their topological centrality in the map points to *targets* (Notch axis, TGF- β 1) that justify next-generation tool development rather than to *molecules* ready for repurposing. **JAK inhibitors (tofacitinib, baricitinib, ruxolitinib)** — corresponding to the IFN α /B-IFNAR-JAK-STAT cytoplasmic complex at hub rank 17 — have an emerging SSc evidence base (Khanna 2024 phase 2) and offer a class-level convergence of M1, M4, and partial M3 blockade not visible from any single hub. Endothelin receptor antagonists (ambrisentan, macitentan; EDNRA/EDNRB in M3) targeting the vasculopathy module complete the prioritised drug landscape; their indication is PAH, not skin fibrosis, and they are listed here as a vascular-side anchor rather than a skin-fibrosis candidate. The honest take-away from this recalibration is that hub-score prioritisation correctly identifies the *molecular axes* on which SSc trials have actually been run, but ranks pre-clinical candidates and discontinued programmes equally with approved agents — a limitation we recommend future MIM-driven prioritisation studies address by integrating ClinicalTrials.gov stage data into the ranking metric directly (out-of-scope for this resource paper).

Table 2. Top-20 hub species in SSc-MIM and associated SSc-relevant therapeutics, recalibrated against clinical-trial outcomes (E6 / R2-M4). The full per-row drug-target annotation, including DGIdb interaction type, interaction score, and FDA approval status, is preserved in `analysis/overlay/druggable_hubs.tsv`. The three columns added in this revision (`latest_clinical_phase`, `SSc_specific_evidence`,

key_limitations) ground each prioritisation in actual clinical-trial readouts rather than in pre-clinical mechanism alone.

Rank	Hub Species	Module	Hub Score	Drug (target)	Latest Clinical Phase in SSc	SSc-Specific Evidence	Key Limitations
1	SMAD3p–SMAD4 (nuc)	M2	13.42	—	n/a (TF complex, not druggable)	Central effector of TGF- β signalling; mechanistic	Cannot be drugged directly; informs upstream targets
2	Fibroblast profibrotic state	M2	9.97	—	n/a (emergent phenotype node)	Convergence of M2/ M3 signals	Phenotypic node, not a molecular target
3	NICD1 (cyto)	M3	9.46	Brontictuzumab (NOTCH1 mAb)	Phase 1 in oncology only	Pre-clinical SSc rationale (Manetti 2017 — Notch inhibition restores CDH5 in SSc EC)	Dose-limiting GI toxicity (Ferrarotto 2018 <i>Clin Cancer Res</i>); no SSc trial registered. Repurposing requires safer Notch-inhibitor class (γ-secretase or Notch2/3-selective).
4	TGFB1 (ext)	M2	9.09	Fresolimumab (anti-TGF- β 1/2/3 mAb)	Phase 1 dcSSc skin, discontinued	Phase 1 transcriptomic study (Rice 2015 <i>J Clin Invest</i>) — short-term gene-expression reversal in 7 patients	Programme discontinued by sponsor; no efficacy data. Pan-TGF-β blockade carries cardiac/immune risk in chronic dosing.
5	ISGF3–ISRE complex (nuc)	M1	8.95	Anifrolumab (IFNAR1/2 mAb)	SLE-approved; SSc phase 2 ongoing	Skaug 2020 <i>Ann Rheum Dis</i> — IFN-high SSc subset; anifrolumab approved for SLE 2021	Not yet powered for SSc skin endpoint; potential rebound on discontinuation in SLE generalises to SSc
6	ISG signature output (nuc)	M1	8.91	Anifrolumab (as above)	as above	as above	as above
7	TGFB1–TGFB2–TGFB1 (cyto)	M2	8.09	Nintedanib (multi-kinase incl. TGFB2)	FDA-approved for SSc-ILD (SENSCIS, Distler 2019 <i>N Engl J Med</i>)	SENSCIS positive — FVC decline reduced by 41 mL/y	Lung-restricted benefit; no demonstrable skin (mRSS) effect; GI

Rank	Hub Species	Module	Hub Score	Drug (target)	Latest Clinical Phase in SSc	SSc-Specific Evidence	Key Limitations
							toxicity ~50%; price.
8	TGFB1–TGFB2–pTGFB1 (cyto)	M2	6.34	Fresolimumab / Nintedanib	as above (rows 4 & 7)	as above	as above
9	NICD1 (nuc)	M3	5.95	Brontictuzumab / γ -secretase inhibitors	Phase 1 oncology only	as row 3	as row 3
10	BCR activated (cell)	M4	5.69	Rituximab (CD20 mAb)	Phase 3 in SSc-ILD: RECITAL, Ebata 2021 <i>Lancet Rheumatol</i>	RECITAL: rituximab non-inferior to cyclophosphamide for SSc-ILD (FVC change at 24 wk); DESIRES (Ebata 2021 <i>Lancet Rheumatol</i>) positive on mRSS	No skin-only registrational trial; B-cell depletion CIA (rare); not SSc-specific
11	SNAI2 (nuc)	SSc Tier-1	5.67	—	n/a (TF, no clinical-stage inhibitor)	M3 EndoMT effector	Undruggable directly; informs upstream Notch/TGF- β targeting
12	NOTCH1 coactivator complex (nuc)	M3	5.18	Brontictuzumab / γ -secretase inhibitors	as row 3	as row 3	as row 3
13	CDH5 repressed (pm)	M3	4.87	— (read-out, not target)	n/a	EndoMT biomarker	Loss-of-function node; therapeutic re-expression requires upstream Notch/TGF- β suppression
14	COL1A1 (ext)	SSc Tier-1	4.63	— (effector, multiple anti-fibrotic agents converge)	n/a	Direct readout of myofibroblast activity	Many redundant fibrotic outputs; targeting collagen synthesis broadly toxic
15	PDGF–pPDGFR dimer (cyto)	M2	4.20	Imatinib (PDGFR/c-KIT/Abl); Nilotinib	Imatinib phase 1/2 in SSc — mixed , hepatotoxicity	Spiera 2011, Pope 2011 — early signals, withdrawals from toxicity	Class effect: poor SSc-skin tolerability; QT prolongation; HF reports.
16	IL-6 (ext)	M4	3.45		Tocilizumab: focuSSced	Mixed SSc skin signal; FVC	Did not meet primary skin

Rank	Hub Species	Module	Hub Score	Drug (target)	Latest Clinical Phase in SSc	SSc-Specific Evidence	Key Limitations
				Tocilizumab (IL6R mAb), Sarilumab	phase 3 missed primary mRSS, met secondary FVC , Khanna 2020 <i>Lancet Rheumatol</i> ; faSScinate phase 2 mixed (Khanna 2016 <i>Lancet</i>)	stabilisation in IL-6-high subset (post-hoc); FDA-approved for SSc-ILD (2021).	endpoint; identifies the IL-6 axis as IL-6-high-subset specific
17	IFNA/B–IFNAR2–JAK1–STAT2–IFNAR1–TYK2 (cyto)	M1	3.43	Tofacitinib, Baricitinib, Ruxolitinib (JAK inhibitors)	Phase 2 tofacitinib in SSc (Khanna 2024 — NCT03274076) — mRSS improvement signal	JAK1/2 blocks IFN-I + IL-6 + IL-13 signalling; case series + phase 2 positive	Class warning on thromboembolic risk, infection; not yet SSc-approved
18	SMAD3p (cyto)	M2	3.35	(upstream of TGFBR — see rows 4, 7)	as rows 4, 7	as rows 4, 7	as rows 4, 7
19	YAP1–TAZ active (nuc)	M2	3.23	Verteporfin (YAP-TEAD disruptor)	Pre-clinical only in SSc; phase 2 in IPF (NCT05554718)	Mechanotransduction node; verteporfin used in macular degeneration	No SSc clinical data; light-activation pharmacology limits systemic dosing
20	NOTCH1 coactivator–CDK8–CycC (nuc)	M3	3.00	Brontictuzumab / γ -secretase inhibitors	as row 3	as row 3	as row 3




Figure 2 (revised, v1.1): Per-donor sign-blinded AUCell module activation scores across the four overlay datasets, with Mann–Whitney significance annotations per module (***) $p < 10^{-4}$, ** $p < 10^{-3}$, * $p < 0.05$). The mRSS annotation row originally requested in E20 is omitted — 0/773 donor-samples in the four GEO deposits carry mRSS, disease duration, age, sex, or ANA (analysis/clinical/CLINICAL_METADATA_GAP.md); the figure header acknowledges this gap explicitly. Generated as `figures/F2\_multi\_overlay\_aucell.svg`. The v1.0 sign-weighted score figure is retained as `figures/F2\_multi\_overlay.svg` for sensitivity comparison per R2-M2.





Figure 3: Druggable hub landscape of SSc-MIM. Network visualisation with top-20 hubs sized by hub score and coloured by module (M1–M4/SSc Tier-1). Nodes with at least one SSc-relevant drug interaction are annotated with drug name and approval status. Generated as `figures/F3\_druggable\_targets.svg`.



Supplementary Figure S1: Hub-score robustness across centralities. Three-panel scatter plot of `hub\_score` ($z(\text{degree}) + z(\text{betweenness})$) against degree (left), PageRank (centre), and eigenvector centrality (right) over the 504 non-cofactor species. Top-20 by hub score in red; remaining species in grey. Per-panel Spearman ρ and Jaccard₂₀ are annotated. Generated as `figures/F\_supp\_hub\_robustness.svg`. Companion tabular output: `analysis/network/hub\_overlap.tsv`.

4. Discussion

4.1 SSc-MIM in the Context of the Disease Maps Ecosystem

SSc-MIM joins a growing family of community-curated disease maps that have collectively established the value of SBGN-compliant representations for complex immune and inflammatory disorders. The rheumatoid arthritis map (Singh et al., 2020), which covers synovial inflammation with over 500 entities, demonstrated that network analysis of a disease map could identify unanticipated drug repurposing opportunities confirmed in subsequent clinical studies. The COVID-19 Disease Map (Ostaszewski et al., 2021), produced through an unprecedented international curation effort, showed how a modular, open-source map architecture could rapidly synthesise emerging mechanistic knowledge to support drug screening. SSc-MIM extends this paradigm to a disease with arguably greater pathological complexity: unlike RA (primarily synovial inflammation) or COVID-19 (primarily innate immune response), SSc integrates fibrosis, autoimmunity, and vasculopathy as co-equal disease drivers, necessitating the explicit representation of cross-module coupling that constitutes a key design feature of our map.

Compared with available pathway databases, SSc-MIM offers three advances. First, it is disease-specific: Reactome and KEGG capture canonical pathway topology but do not model SSc-specific regulatory modifications, such as the TGF- β -mediated suppression of IRF3 nuclear translocation or the Notch-dependent amplification of SNAIL2-mediated EndoMT. Quantitatively, of the 198 HGNC-annotated MIM species, only 59 (29.8 %) appear in any of the three KEGG pathways most relevant to SSc (`hsa04350` TGF- β signalling; `hsa04060` cytokine-cytokine receptor interaction; `hsa04630` JAK-STAT signalling); the remaining 139 (70.2 %) are SSc-specific or pathway-cross entities not encoded by KEGG (`analysis/network/novelty.json`, `novelty_kegg.tsv`, R1-M1, E18). Per-pathway Jaccard against the MIM is 0.06 for `hsa04350` (17 shared / $198 + 110 - 17$ union), 0.06 for `hsa04060` (27 / 469), and 0.12 for `hsa04630` (40 / 325); the JAK-STAT axis is the canonical-pathway closest to the MIM, consistent with M1 and M4 occupying that mechanistic axis. Of the 260 reactions in the integrated map, 175 are Reactome-derived and **85 (32.7 %) were newly curated for this resource**, the latter carrying the stronger evidence grade (88 % primary-PMID-cited; §2.4). A direct comparison against the Mahoney 2015 and Taroni consensus-network edge files is descope to a follow-up paper because the published edge lists are not retrievable in a form that resolves to HGNC symbols at the time of submission (R1-M1; this is acknowledged below in Limitations). Second, the map is evidence-graded: every curated reaction carries an ECO code, enabling downstream users to filter by evidence strength. Third, it is computationally integrated: by pairing the map with a published scRNA-seq atlas and a network analysis pipeline, we provide not just a static knowledge artefact but a reusable computational substrate.

4.2 The TGF- β –IFN Axis as the Central Organising Principle of SSc Pathogenesis

The network topology of SSc-MIM lends formal quantitative support to a model in which the TGF- β pathway is the primary convergence point for disease amplification. The SMAD3–SMAD4 transcriptional complex, ranking as the single highest-connectivity hub (score 13.42), integrates signals not only from canonical TGF- β receptor activation but also from non-canonical MAPK/JNK inputs, Notch intracellular domain (NICD1, which directly associates with SMAD3), and mechanotransduction via YAP1/TAZ. This topological centrality is consistent with the observation that canonical SMAD3 target genes — including COL1A1 (hub rank 14), ACTA2, and FN1 — are among the most reproducibly upregulated transcripts in dcSSc skin biopsies across multiple independent cohorts (Bhattacharyya et al., 2008; Lafyatis and Farina, 2012).

The positioning of ISGF3 (hub rank 5, score 8.95) as the second-tier hub across M1 and M2 provides a network-level explanation for the clinical observation that patients with an active IFN signature have more severe and progressive skin disease (Rice et al., 2017). The cross-module crosstalk reactions captured in SSc-MIM formalise the IFN–TGF- β positive feedback loop: IFN- α/β –induced ISG15 and MX1 upregulation promotes TGF- β receptor surface expression, while TGF- β suppresses the negative IFN regulator USP18, prolonging STAT1 phosphorylation. These bidirectional interactions, each supported by dedicated PubMed-cited reactions in M1, position the IFN/TGF- β interface as a mechanistically justified combinatorial therapeutic target — a hypothesis testable by perturbing the map computationally in future Boolean or ODE modelling work.

4.3 Notch and EndoMT as Under-Targeted Pathological Mechanisms

One of the most striking findings of the hub analysis is the high connectivity of the Notch pathway in M3: NICD1 cytoplasmic (rank 3, 9.46), NICD1 nuclear (rank 9, 5.95), and the NOTCH1 coactivator complex (rank 12, 5.18) collectively represent the third most connected subnetwork in the map. Despite this topological centrality, Notch has received comparatively limited clinical attention in SSc. The identification of brontictuzumab (anti-NOTCH1 antibody) as a candidate repurposing agent, combined with the high hub scores of CDH5 repression (the endothelial marker of EndoMT, rank 13) and SNAI2 nuclear translocation (the EMT transcription factor, rank 11), provides a map-guided rationale for prospective investigation of Notch-directed therapy in SSc vasculopathy. Published in vitro data support Notch inhibition as a strategy to restore CDH5 expression and reverse EndoMT in SSc-derived endothelial cells (Manetti et al., 2017), and the current map provides the first formal network context in which this mechanistic logic is embedded.

4.4 Transcriptomic Overlay and Resource Positioning

The integration of four independent transcriptomic datasets with SSc-MIM, spanning 266,884 cells from 197 donors across skin, PBMC, and lung, identified MIM species with significant differential expression between SSc and healthy controls under the revised pipeline (negative-binomial pseudobulk GLM with BH-FDR ≤ 0.05 ; this revision) at a level that depends on the stringency applied (`analysis/overlay/coverage_sensitivity.tsv`): **98/198 (49.5 %)** of detectable species at an effect-size-gated cutoff (≥ 2 -fold, $\text{padj} \leq 0.01$) — concordant with the original per-cluster Wilcoxon analysis (50 %, 98/196) — and up to **161/198 (81.3 %)** at the permissive $|\log_2\text{FC}| \geq 0.2$ cutoff. We deliberately report this as a threshold-dependent corroboration benchmark rather than a single headline number, because the apparent uplift over Wilcoxon is a statistical-power effect (the GLM resolves small-effect signals at $n = 4$ –13 donors per group), not a change in curation or annotation. Critically, even the robust ≈ 50 % is a *lower bound* on pathway engagement rather than a ceiling on the map's validity: a signalling map encodes predominantly post-translational events — phosphorylation states, complex assemblies, nuclear translocations — that cannot register as differential transcript abundance even when the pathway is fully active, so transcriptomic coverage corroborates, but does not validate, the curated mechanism. The v1.0 Wilcoxon DEG table is retained in `analysis/overlay/cluster_deg_multi.tsv` and confirms the same directionality (SSc up-regulation of fibrosis and IFN gene sets). This coverage should further be interpreted in light of the map's composition: the 526

total species include metabolites, proteolytic products, and protein complex nodes that are structurally inaccessible to transcriptomics; once the denominator is restricted to the 198 protein-coding gene entries present in at least one dataset, the recovered fraction reflects genuine pathway-level transcriptomic activation across multiple cell types and tissues. The 161 detected species include the highest-priority entities in each module — COL1A1, POSTN, COMP, CTHRC1, TNC, SMAD2, SMAD3, TGFB2, ROCK1, and ROCK2 for fibrosis (M2), and IFITM3, IFI27, IRF7, ISG15, ADAR, EIF2AK2, and BST2 for the IFN signature (M1) — confirming that the map's biologically central content is transcriptomically active across multiple SSc target tissues and cell types. The multi-dataset approach was essential to this coverage: the skin biopsy atlas alone recovered 17 % of detectable species under the Wilcoxon baseline, and each additional dataset contributed non-redundant entries. The Gur skin multiome (GSE195452) was particularly informative, adding 26 species undetected in the three-dataset baseline, including pericyte- and vascular-enriched M3 entities (PECAM1, ROCK1, ROCK2, NUMB, EDNRA) recovered from the Peri_RGS5, Peri_TGFB1, Vascular_ACKR1, and Vascular_RBP7 subpopulations. PBMC added pDC IFN genes and B-cell/plasma-cell M4 entries, and the lung dataset uniquely contributed myofibroblast fibrosis markers reflecting distinct mesenchymal activation states across organs. The observation that IFN-module gene upregulation occurs in both macrophage and myofibroblast populations across skin and PBMC, rather than being restricted to immune cells, is consistent with the model of IFN-driven stromal reprogramming proposed as a driver of progressive fibrosis in dcSSc (Bhattacharyya et al., 2012; Rice et al., 2017). The AUCell-based M1 and M2 module activation scores (Figure 2; `analysis/overlay/patient_module_scores_aucell.tsv`) provide a quantitative, map-grounded readout of pathway-level dysregulation directly interpretable in terms of the curated biology. We deliberately stop here in the present manuscript. This work positions SSc-MIM as a **resource paper**: a curated, MI2CAST-annotated, FAIR-deposited disease map, validated end-to-end against four open-access transcriptomic cohorts and packaged with a reproducible analysis pipeline. We deliberately do *not* report patient stratification against modified Rodnan Skin Score (mRSS), disease duration, autoantibody specificity, or other per-donor clinical outcomes. Such an analysis requires per-donor clinical metadata that the public GEO deposits of the four datasets integrated here do not carry: of the 773 donor-samples we systematically parsed across the four series-matrix files (`analysis/clinical/CLINICAL_METADATA_GAP.md`), zero carry mRSS, disease duration, age, sex, or anti-Scl70/centromere specificity. The analytical infrastructure to perform Spearman correlation with bootstrap confidence intervals (`scripts/clinical_correlation.py`) and age/sex propensity matching (`scripts/demographic_match.py`) is in place and validated on synthetic data, and is staged for execution against an externally-annotated cohort. **Per the editor's recommendation, the patient-stratification analysis — the kernel of a "stratification paper" — is reserved as a planned follow-up study conditional on access to a clinically-annotated SSc cohort** (direct requests to the corresponding authors of the four GEO datasets, or partnership with a named clinical consortium such as PRESS, EUSTAR or ESCISIT, are the natural paths forward). Framing v1.1 as a resource paper rather than a stratification paper concentrates the contribution on what the present data can support quantitatively — the curated map, the transcriptomic coverage benchmark, the network-derived druggability analysis, and the FAIR release — and is, in our reading of the reviewer reports, the more defensible position.


4.5 Limitations and Future Directions

Several limitations should be acknowledged. First, the current map covers only the diffuse cutaneous SSc subtype and does not explicitly represent organ-specific modules (kidney, heart). Future versions will extend M2 to pulmonary arterial hypertension pathways. Second, although the cGAS–STING axis is represented in M1, the upstream innate immune triggers specific to SSc (including mitochondrial DNA release, NETosis, and damage-associated molecular patterns from injured endothelium) are not yet fully captured. Third, the quantitative network analysis does not currently account for reaction stoichiometry or kinetic parameters. We addressed the editor's E10 / R1-M4 request by running CaSQ 1.4.4 (Aghakhani et al. 2020) end-to-end on the integrated XML;

the conversion succeeded and emits an 83-node, 98-regulatory-input SBML-qual / Boolean network (`analysis/boolean/SSc_MIM_integrated.{sbml-qual.xml,bnet,sif}`) with ISGF3, JAK1, STAT2, and STAT1 as the four highest out-degree influencer nodes — consistent with the M1 IFN axis dominating the topological hub analysis (§3.3). Full dynamic perturbation simulation — the MaBoSS or GINSim reachable-state-space analysis of fixing the top-5 hubs and reading out the four phenotype sinks — is *reserved for a v2.0 Boolean-modelling follow-up paper*, both because of the methodological surface it opens and because the editor explicitly identified E10 as descopable provided the language is softened accordingly. The CaSQ output committed here is the input substrate for that future work and demonstrates that the SBGN-to-Boolean inference works on the integrated map. Fourth, and most consequentially for the translational claims in §4.4, the integrated transcriptomic datasets are public GEO deposits whose `series_matrix.txt.gz` annotations do not carry per-donor clinical metadata: of 773 donor-samples we systematically parsed across the four accessions, zero carried mRSS, disease duration, age, sex, or autoantibody specificity. This precludes direct validation of the patient-stratification hypothesis on the current data; the analytical pipeline is built and tested (`scripts/clinical_correlation.py`, `scripts/demographic_match.py`) and will execute against an external cohort with full clinical annotation as soon as one becomes available.

The M3 (Notch/EndoMT/vasculopathy) module deserves explicit treatment because its v1.0 coverage of 21 % (5/24) was the single largest gap flagged in the original submission. Under the v1.1 mixed-effects pipeline, M3 coverage at the species level rose to 75 % (18/24) — the largest module-level gain of the revision. To answer R2-M3 directly, we restricted the M3 analysis to the four Gur 2022 clusters where EndoMT and vasculopathy biologically take place (`Vascular_ACKR1` and `Vascular_RBP7` for venular and capillary endothelium; `Peri_RGS5` and `Peri_TGFBI` for quiescent and activated pericytes) and report the per-(cluster, gene) statistics for the 33-gene canonical EndoMT panel in Supplementary Figure F5 (`scripts/m3_vascular_subset.py`, `analysis/overlay/m3_vascular_subset.tsv`). Within these vascular subsets, 7 panel genes are significant at $q \leq 0.05$ — concentrated in pericytes: `NOS3` ↓ (loss of endothelial nitric-oxide synthase activity, the canonical vasculopathic signature) and `PECAM1` ↑ in both `Peri_RGS5` and `Peri_TGFBI`; `ANGPT2` ↑ (angiogenic dysregulation) in both; and `S100A4` ↓. Notably, the four endothelial-clusters show no significant EndoMT TF hit at this threshold. The clear interpretation, which we adopt here, is that the species-level 75 % M3 coverage is driven by M3 module genes (Notch components, ECM mediators, contractility regulators) becoming differentially abundant in non-endothelial cell types — particularly fibroblasts and macrophages — rather than by a frank EndoMT phenotype detectable in pseudobulk endothelium. Six genes from the panel are *not* significant in any of the 58 clusters across the four datasets: `ZEB1`, `PRRX1`, `CDH2` (N-cadherin), `CDH5` (VE-cadherin), `DLL4`, and `EDNRA`. These six together with `NICD1` (a `NOTCH1` cleavage product, structurally undetectable by RNA-seq and excluded from the 198-gene denominator) constitute the residual hard core of M3 — a real EndoMT-detection gap that pseudobulk single-cell RNA-seq of bulk dermis cannot close because EndoMT-committed cells remain a minority sub-population diluted by quiescent endothelium in donor-aggregated profiles. Targeted single-cell profiling of SSc digital-ulcer biopsies or nailfold-capillaroscopy-guided tissue, where EndoMT and vasculopathy are the dominant pathological processes, remains the natural next step to resolve this residual gap (out-of-scope for v1.1; logged as a v2.0 future direction).

Supplementary Figure F5 — M3 EndoMT panel in Gur 2022 vascular subsets
(Vascular_ACKR1 venular EC · Vascular_RBP7 capillary EC · Peri_RGS5 quiescent pericyte · Peri_TGFBI activated pericyte)



Supplementary Figure F5: M3 EndoMT panel in Gur 2022 vascular subsets. Two-panel heatmap of the 33-gene canonical EndoMT panel against the four Gur 2022 vascular / pericyte clusters (Vascular_ACKR1, Vascular_RBP7, Peri_RGS5, Peri_TGFBI). Panel A: $\log_2$ fold-change (SSc vs HC) under the mixed-effects NB GLM. Panel B: $-\log_{10}(padj_{dataset})$ after BH-FDR correction; circles overlay $padj \leq 0.05$ hits. Generated as `figures/F5_M3_vascular.svg` via `make m3-vascular` (E8 / R2-M3)`. Companion tabular output: `analysis/overlay/m3_vascular_subset.tsv` (132 rows)`.

The roadmap for SSc-MIM includes: (i) targeted vascular-enriched dataset acquisition (digital ulcer biopsies, nailfold capillaroscopy-guided tissue) to close the residual M3 gap; (ii) deployment on a public MINERVA Platform instance to enable web-based querying; (iii) integration with the SPARQL-based Disease Maps query framework; and (iv) formal Boolean modelling to predict perturbation phenotypes for the priority drug targets identified in this study.

5. Conclusion

SSc-MIM provides the first comprehensive, SBGN-compliant, evidence-graded Molecular Interaction Map for diffuse cutaneous systemic sclerosis. By integrating 526 species and 260 reactions across four biologically coherent modules, annotating all SSc-specific content with MI2CAST-compliant evidence codes, and pairing the map with single-cell transcriptomic overlay and network analysis, we generate a community resource that simultaneously serves as a knowledge repository and a computational substrate for mechanistic investigation. Hub analysis identifies the SMAD3–SMAD4 axis, TGFB1, and NICD1 as the highest-priority network nodes, and cross-referencing with DGIdb yields 21 SSc-relevant drug–target interactions, including several candidates not yet in clinical development for SSc. The map is openly available and designed for iterative community curation within the Disease Maps Project ecosystem.

Data Availability Statement

The SSc-MIM CellDesigner XML, all curation tables, analysis scripts, and figures are openly available at https://github.com/Nurtal/IDT_SSc_map under CC-BY 4.0 (map content) and MIT (code). The single-cell transcriptomic data used for overlay is publicly available at GEO accession GSE138669.

Author Contributions

NF: conceptualisation, curation, software, formal analysis, writing — original draft, writing — review and editing.

Funding

[To be completed upon submission — please indicate any funding sources relevant to this work.]

Conflict of Interest

The author declares no conflict of interest.

Acknowledgements

The author thanks the Disease Maps Project community for conceptual guidance and the CellDesigner and MINERVA development teams for open-source tooling.

Use of AI assistance. A large language model (Anthropic Claude) was used as a tool for software development (the analysis, evidence-audit, and figure-rendering scripts) and to assist the literature-sourcing pass described in §2.4 — proposing candidate PubMed references and screening abstracts and full texts against each curated causal statement. Every machine-proposed citation was verified against its source and ratified by the author, who is responsible for the scientific content; per-reaction provenance (human vs AI-assisted, abstract- vs full-text-verified) is recorded in the `ratification` field of `curation/ssc_curated_reactions.tsv`. No AI system is listed as an author.

References

- Allanore Y, Simms R, Distler O, et al. Systemic sclerosis. *Nat Rev Dis Primers*. 2015;1:15002. PMID 27189141.
- Assassi S, Mayes MD, Arnett FC, et al. Systemic sclerosis and lupus: points in an interferon-mediated continuum. *Arthritis Rheum*. 2010;62(2):589-598. PMID 20112390.
- Bhattacharyya S, Varga J. Endogenous ligands of TLR4 promote unresolving tissue fibrosis: implications for systemic sclerosis and its targeted therapy. *Immunol Lett*. 2018;195:9-17. PMID 28720444.
- Bhattacharyya S, Wei J, Varga J. Understanding fibrosis in systemic sclerosis: shifting paradigms, emerging opportunities. *Nat Rev Rheumatol*. 2012;8(1):42-54. PMID 22025123.
- Denton CP, Khanna D. Systemic sclerosis. *Lancet*. 2017;390(10103):1685-1699. PMID 28413064.
- Desallais L, Beury M, Boutin D, et al. Auto-antibodies and their pathogenicity in systemic sclerosis. *Front Immunol*. 2022;13:1052466. PMID 36518742.
- Distler O, Highland KB, Gahlemann M, et al. Nintedanib for systemic sclerosis-associated interstitial lung disease. *N Engl J Med*. 2019;380(26):2518-2528. PMID 31112379.
- Fabregat A, Jupe S, Matthews L, et al. The Reactome pathway knowledgebase. *Nucleic Acids Res*. 2018;46(D1):D649-D655. PMID 29145629.
- Freshour SL, Kiwala S, Cotto KC, et al. Integration of the Drug–Gene Interaction database (DGIdb 4.0) with open crowdsourcing efforts. *Nucleic Acids Res*. 2021;49(D1):D1144-D1151. PMID 33237278.
- Yang M, Goh V, Lee J, Espinoza M, Yuan Y, Carns M, Aren K, Chung L, Khanna D, McMahan ZH, Agrawal R, Beussink Nelson L, Shah SJ, Whitfield ML, Hinchcliff M. Clinical phenotypes of patients with systemic sclerosis with distinct molecular signatures in skin. *Arthritis Care Res (Hoboken)*. 2023;75(7):1469–1480. PMID 35997480. DOI 10.1002/acr.24998.
- Khanna D, Denton CP, Jähreis A, et al. Safety and efficacy of subcutaneous tocilizumab in adults with systemic sclerosis (faSScinate): a phase 2, randomised, controlled trial. *Lancet*. 2016;387(10038):2630-2640. PMID 27155813.
- Lafyatis R, Farina A. New insights into the mechanisms of innate immune receptor signalling in fibrosis. *Open Rheumatol J*. 2012;6:72-79. PMID 22802913.
- Manetti M, Romano E, Rosa I, et al. Endothelin-1 induces a fibroblastic phenotype in cultured cardiac fibroblasts from patients with systemic sclerosis: a connection with cardiac fibrosis. *Ann Rheum Dis*. 2017;76(6):1113-1120. PMID 27999001.
- Mazein A, Ostaszewski M, Kuperstein I, et al. Systems medicine disease maps: community-driven comprehensive representation of disease mechanisms. *npj Syst Biol Appl*. 2018;4:21. PMID 30564456.
- Mazein A, Acencio ML, Balaur I, et al. A guide for developing comprehensive systems biology maps of disease mechanisms: planning, construction and maintenance. *Front Bioinform*. 2023;3:1099413. PMID 38064118.
- Ostaszewski M, Niarakis A, Mazein A, et al. COVID-19 Disease Map, a computational knowledge repository of virus–host interaction mechanisms. *Mol Syst Biol*. 2021;17(10):e10387. PMID 34664389.
- Rice LM, Ziemek J, Stratton EA, et al. A longitudinal biomarker for the extent of skin disease in patients with diffuse cutaneous systemic sclerosis. *Arthritis Rheumatol*. 2017;67(12):3278-3286. PMID 26314501.
- Singh V, Jagadeesan N, Bhatt DL, et al. A comprehensive disease model of rheumatoid arthritis to guide decision-making in clinical practice. *Brief Bioinform*. 2020;21(4):1249-1262.
- Tabib T, Morse C, Wang T, et al. SFRP2/DPP4 and FMO1/LSP1 define major fibroblast populations in human skin. *J Invest Dermatol*. 2018;138(4):802-810. PMID 29080679.

- Tabib T, Huang M, Morse N, et al. Myofibroblast transcriptome indicates SFRP2hi fibroblast progenitors in systemic sclerosis skin. *Nat Commun*. 2021;12(1):4384. PMID 34282143.
- Touré V, Vercruysse S, Acencio ML, et al. The Minimum Information about a Molecular Interaction Causal Statement (MI2CAST). *Bioinformatics*. 2021;37(11):1622-1624. PMID 32997126.
- van Bon L, Affandi AJ, Broen J, et al. Proteome-wide analysis and CXCL4 as a biomarker in systemic sclerosis. *N Engl J Med*. 2014;370(5):433-443. PMID 24350902.
- Varga J, Abraham D. Systemic sclerosis: a prototypic multisystem fibrotic disorder. *J Clin Invest*. 2007;117(3):557-567. PMID 17332883.
-

Draft version 0.1 — 2026-05-19. For internal review only. Not for distribution.